

An end-to-end ML project

Mauricio A. Álvarez

Machine Learning and Adaptive Intelligence
The University of Sheffield



The
University
Of
Sheffield.

Overview

- This week's session will show an example of an end-to-end project in ML. The focus is on supervised learning.

- The lecture provides a general overview and the Jupyter Notebook will provide a hands-on exercise.

- We will describe the typical process involved in providing an ML solution.

- We will use *scikit-learn* (<https://scikit-learn.org/stable/>) to illustrate several of the steps involved.

Why *scikit-learn*?

- scikit-learn is a machine learning library for Python.
- It supports supervised and unsupervised learning.
- It provides several utilities for data preprocessing and model selection and evaluation.
- More importantly: it is open source (BSD licence).

Machine learning project checklist

1. Frame the problem and look at the big picture.
2. Get the data.
3. Explore the data to get insights.
4. Prepare the data to better expose the underlying data patterns.
5. Explore many different models and shortlist the best ones.
6. Fine-tune your models and combine them into a solution.
7. Present your solution.

Contents

Frame the problem and look at the big picture

Get the data

Explore the data

Prepare the data

Shortlist models and then fine-tune them

Present your solution

Set the problem in context (I)

- How will your solution be used?
- What are the current solutions to the problem (if any)?
- Is it a regression or a classification problem?
- How is performance measured?

Set the problem in context (II)

- What would be the minimum performance necessary?
- Are there comparable problems?
- Is human expertise available?

Model performance assessment

- In ML, we use several metrics to assess the performance of a model.
- Performance measures commonly used in regression are
 - root mean squared error (RMSE).
 - absolute mean error (MAE).
- For classification, metrics of performance include:
 - confusion matrix.
 - precision/recall.
 - accuracy.

Metrics in regression

- We have a predictive model $f(\mathbf{x}, \mathbf{w})$ and we want to assess its performance over a dataset of instances $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)$.
- The root mean squared error is defined as

$$\text{RMSE} = \sqrt{\frac{1}{N} \sum_{i=1}^N [y_i - f(\mathbf{x}_i, \mathbf{w})]^2}.$$

- The mean absolute error is defined as

$$\text{MAE} = \frac{1}{N} \sum_{i=1}^N |y_i - f(\mathbf{x}_i, \mathbf{w})|.$$

Metrics in classification: confusion matrix

- We have a predictive model $f(\mathbf{x}, \mathbf{w})$ and we want to assess its performance over a dataset of instances $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)$.
- Binary classification problem: spam detection. There are two classes, spam and not-spam.
- The confusion matrix is a table that summarizes how successful the classification model is at predicting examples belonging to various classes

	spam (predicted)	not-spam (predicted)
spam (actual)	23 (TP)	1 (FN)
not-spam (actual)	12 (FP)	556 (TN)

where TP stands for **true positive**, TN stands for **true negative**, FP stands for **false positive** and FN stands for **false negative**.

Metrics in classification: precision/recall (I)

- **Precision** is the ratio of correct positive predictions to the overall number of positive predictions

$$\text{precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}.$$

- **Recall** is the ratio of correct positive predictions to the overall number of positive examples in the dataset

$$\text{recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}.$$

Metrics in classification: precision/recall (II)

- ❑ Say you look for the terms “pattern recognition and machine learning” in the StarPlus University’s Library website.
- ❑ Precision is the fraction of relevant documents in the list of returned documents (“how valid the results are?”)
- ❑ Recall is the fraction of relevant documents returned to the total number of relevant documents that could have been returned (“how complete the results are?”)
- ❑ In the spam example, we would like high precision, this is we would like the FP to be low meaning having few genuine messages predicted as spam.
- ❑ We are probably Ok with a low recall, some spam messages sent to Inbox because they are predicted as not-spam.

Metrics in classification: accuracy

- **Accuracy** is the ratio of examples correctly classified over the total number of examples classified

$$\text{accuracy} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}.$$

- Accuracy is a useful metric when errors predicting all classes are equally important.
- In the spam example, FP are worse than FN.
- Accuracy can be misleading in imbalance classification problems, e.g. you can have a high a TP value and a low TN value, and your accuracy could still be high.

Contents

Frame the problem and look at the big picture

Get the data

Explore the data

Prepare the data

Shortlist models and then fine-tune them

Present your solution

Be aware of

- what data do you need? is it enough?
- legal obligations regarding the data. Who owns the data? Do you have to sign an NDA to work with the data? Do you require a legal agreement?
- For example, since 2018 the UK abides to the General Data Protection Regulation (GDPR) under EU law.
- removing sensitive information (e.g. anonymisation)
- institutions usually have in place a data ethics advisory committees.
- check the size and type of the data (e.g time series? images? spatio-temporal?)
- sample a test set and do not look at it (e.g. avoid data snooping).

Where do we get data from?

UC Irvine Machine Learning Repository


Machine Learning Repository
Center for Machine Learning and Intelligent Systems

About Citation Policy Donate a Data Set Contact
Repository Web Google+
[View ALL Data Sets](#)

Welcome to the UC Irvine Machine Learning Repository!

We currently maintain 557 data sets as a service to the machine learning community. You may [view all data sets](#) through our searchable interface. For a general overview of the Repository, please visit our [About](#) page. For information about citing data sets in publications, please read our [citation policy](#). If you wish to donate a data set, please consult our [donation policy](#). For any other questions, feel free to [contact the Repository librarians](#).

Supported By:  In Collaboration With:  Rexa.info

Latest News:	Newest Data Sets:	Most Popular Data Sets (hits since 2007):
<p>09-24-2018: Welcome to the new Repository admins Dheeru Dua and Efi Karra Taniskidou!</p> <p>04-04-2013: Welcome to the new Repository admins Kevin Bache and Moshe Lichman!</p> <p>03-01-2010: Note from donor regarding Netflix data</p> <p>10-16-2009: Two new data sets have been added.</p> <p>09-14-2009: Several data sets have been added.</p> <p>03-24-2008: New data sets have been added!</p> <p>06-25-2007: Two new data sets have been added: UCI Pen Characters, MAGIC Gamma Telescope</p>	<p>07-22-2020:  Facebook Large Page-Page Network</p> <p>07-17-2020:  Amphibians</p> <p>07-12-2020:  Early stage diabetes risk prediction dataset</p> <p>06-28-2020:  Taiwanese Bankruptcy Prediction</p> <p>06-20-2020:  South German Credit (UPDATE)</p>	<p>3551844:  Iris</p> <p>1931560:  Adult</p> <p>1490003:  Wine</p> <p>1331678:  Breast Cancer Wisconsin (Diagnostic)</p> <p>1315682:  Heart Disease</p>

Featured Data Set: Forest Fires

Task: Regression
Data Type: Multivariate
Attributes: 13

<https://archive.ics.uci.edu/ml/index.php>

Where do we get data from?


Kaggle

Start with more than a blinking cursor

Kaggle offers a no-setup, customizable, Jupyter Notebooks environment. Access free GPUs and a huge repository of community published data & code.

 REGISTER WITH GOOGLE

Register with Email



The screenshot shows the Kaggle Jupyter Notebook interface. The main area displays Python code for training an XGBoost model on a dataset of malicious websites. The sidebar includes sections for Sessions (showing a draft session using GPU off), Workspace (with an input folder), Versions, and Settings (language set to Python, Docker latest available, GPU on, Internet off). The notebook title is "Predict Malicious Websites: XGBoost".

```
# Clean up column names
data.columns = [col.lower().strip() for col in data.columns]

# remove non-numeric columns
data = data.select_dtypes(['number'])

# split data into training & testing
train, test = train_test_split(data, shuffle=True)

# peek @ dataframe
train.head()

# split training data into inputs & outputs
X = train.drop(['type'], axis=1)
y = train['type']

# specify model (xgboost defaults are generally fine)
model = xgb.XGBRegressor(tree_method = 'gpu_hist')

# fit our model
model.fit(y=y, X=X)

# split testing data into inputs & output
test_X = test.drop(['type'], axis=1)
test_y = test['type']
```

<https://www.kaggle.com/>

Where do we get data from?

Amazon's AWS datasets

Registry of Open Data on AWS



About

This registry exists to help people discover and share datasets that are available via AWS resources. Learn more about sharing data on AWS.

See all usage examples for datasets listed in this registry.

See datasets from Facebook Data for Good, NASA Space Act Agreement, NIH STRIDES, NOAA Big Data Project, Space Telescope Science Institute, and Amazon Sustainability Data Initiative.

Search datasets (currently 188 matching datasets)

Search datasets

Add to this registry

If you want to add a dataset or example of how to use a dataset to this registry, please follow the instructions on the [Registry of Open Data on AWS GitHub repository](#).

Unless specifically stated in the applicable dataset documentation, datasets available through the Registry of Open Data on AWS are not provided and maintained by AWS. Datasets are provided and maintained by a variety of third parties under a variety of licenses. Please check dataset licenses and related documentation to determine if a dataset may be used for your application.

The Cancer Genome Atlas

[cancer](#) [genomic](#) [life sciences](#) [STRIDES](#) [whole genome sequencing](#)

The Cancer Genome Atlas (TCGA), a collaboration between the National Cancer Institute (NCI) and National Human Genome Research Institute (NHGRI), aims to generate comprehensive, multi-dimensional maps of the key genomic changes in major types and subtypes of cancer. TCGA has analyzed matched tumor and normal tissues from 11,000 patients, allowing for the comprehensive characterization of 33 cancer types and subtypes, including 10 rare cancers. The dataset contains open Clinical Supplement, Biospecimen Supplement, RNA-Seq Gene Expression Quantification, miRNA-Seq Isoform Expression Quantificati...

[Details →](#)

Usage examples

- [GDC Legacy Archive by National Cancer Institute](#)
- [Comprehensive Characterization of Cancer Driver Genes and Mutations by Matthew H. Bailey, Collin Tokheim, et al.](#)
- [Genomic and Functional Approaches to Understanding Cancer Aneuploidy by Alison M. Taylor, Juliann Shih, et al.](#)
- ["Before and After: A Comparison of Legacy and Harmonized TCGA Data at the Genomic Data Commons" by Galen F. Gao, Joel S. Parker, et al.](#)
- [Using TCGA Data, Resources, and Materials by National Cancer Institute](#)

[See 29 usage examples →](#)

<https://registry.opendata.aws/>

Where do we get data from?

- ❑ Meta portals listing open data repositories.
- ❑ Data portals (<http://dataportals.org/>)
- ❑ Open data monitor (<https://www.opendatemonitor.eu/>)
- ❑ Quandl (<https://www.quandl.com/>)

Three sets: training, validation and test sets

- ❑ In the previous session, we mentioned that we use a dataset to train an ML model.
- ❑ Actually, when we get the data, we should partition the data into three sets
 - training set.
 - validation set.
 - test set.
- ❑ The partition of the dataset into these three sets is done randomly.
- ❑ The training set is usually the biggest one.
- ❑ The validation and test sets have roughly the same amount of samples.

Three sets: what for?

- The training set is used to fit the model using the objective function.
- The validation set is used to choose the best predictive model among a set of candidates.
- The test set is used to assess the final performance of the model before shipping the model to production.
- You can use the training and validation sets as you prefer.
- As we said before, never look at your test set when designing your model, only until you have decided what model to use based on the validation set.

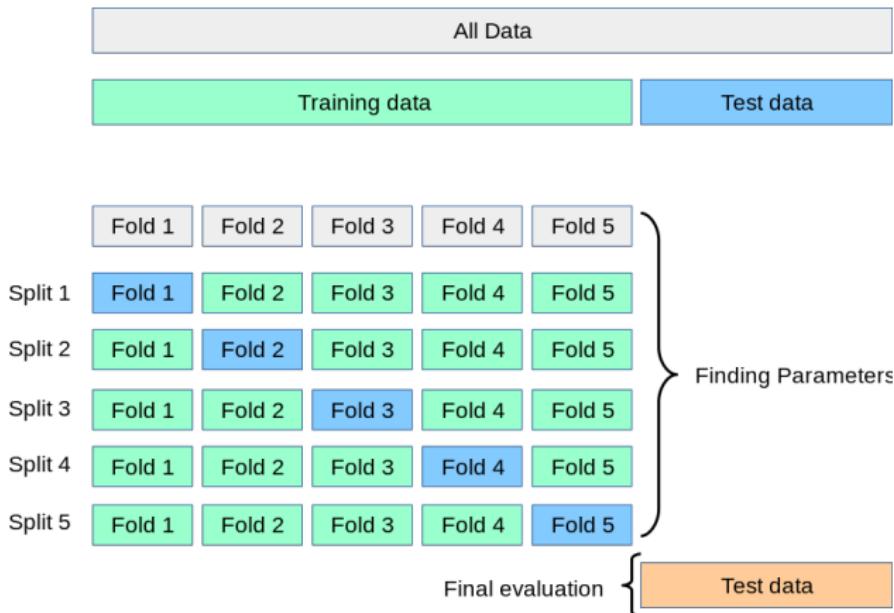
Three sets: large datasets

- ❑ When you have a large dataset, you can have 70% data for training, 15% for validation and 15% for testing.
- ❑ Nowdays, you can have a dataset with millions of data instances. In those cases, you might want to try 95% for training, 2.5% for validation and 2.5% for testing.

Three sets: smaller datasets

- ❑ If your data is small, say in the hundreds or just a few thousands, you want to use as many training instances as possible.
- ❑ We can use a strategy known as *cross-validation* to make better use of the training/validation sets.

k -fold cross-validation



The extreme case is when $k = N$, the number of folds is equal to the number of instances in the training/validation set: leave-one-out (LOO) cross-validation.

Contents

Frame the problem and look at the big picture

Get the data

Explore the data

Prepare the data

Shortlist models and then fine-tune them

Present your solution

To take into account (I)

- Create a copy of the data for exploration.
- If you have a large dataset, you can sample a fraction of the data for exploring it.
- Create a Jupyter Notebook so that you keep track of your data exploration process.
- Study each feature (inputs and outputs) and its characteristics including
 - name
 - type (continuous, bounded/unbounded, categorical, etc.)
 - percentage of missing values
 - noisiness and type of noise (e.g. outliers)
 - Is this feature useful for the task?
 - type of distribution (Gaussian, Gamma, logarithmic, etc.)

To take into account (II)

- Identify the target attribute (for supervised learning).
- Visualise the data using histograms, scatterplots, maps, etc.
- Study correlations between attributes.
- Identify what transformations you might be able to apply.
- Is there any additional data that you might find useful? Can you get it?

Predicting bike rentals




By Sarah Stierch - Own work, CC BY 4.0, <https://commons.wikimedia.org/w/index.php?curid=16075615>

Dataset

- ❑ The feature vector \mathbf{x} includes the following features: hour, temperature, humidity, wind speed, visibility, Dew point temperature, solar radiation, rainfall, snowfall, seasons, holiday, functioning day.
- ❑ The variables hour, temperature, humidity, wind speed, visibility, Dew point temperature, solar radiation, rainfall, snowfall can be considered as continuous.
- ❑ The variables seasons, holiday, and functioning day are categorical variables.
- ❑ The output variable y is the number of bikes rented.

Scatter plot



Study correlations between attributes

The correlation coefficient between two RVs X and Y is given as


$$\rho_{X,Y} = \frac{E\{(X - \mu_X)(Y - \mu_Y)\}}{\sigma_X \sigma_Y} = \frac{\sigma_{X,Y}}{\sigma_X \sigma_Y},$$

where $-1 < \rho_{X,Y} < 1$ and $\sigma_{X,Y}$ is known as the covariance between X and Y .

```
In [18]: corr_matrix["Rented Bikes"].sort_values(ascending=False)
```

```
Out[18]: Rented Bikes      1.000000
Temperature        0.538558
Hour              0.410257
Dew point temperature  0.379788
Solar Radiation    0.261837
Visibility         0.199280
Wind speed         0.121108
Rainfall           -0.123074
Snowfall           -0.141804
Humidity          -0.199780
Name: Rented Bikes, dtype: float64
```

Correlation and dependence



By DenisBoigelot, original uploader was Imagecreator - Own work, original uploader was Imagecreator, CC0,

<https://commons.wikimedia.org/w/index.php?curid=15165296>

Contents

Frame the problem and look at the big picture

Get the data

Explore the data

Prepare the data

Shortlist models and then fine-tune them

Present your solution

Data cleaning (I)

- ❑ Remove the outliers in your data (optional).
- ❑ Handle the missing values
 - filling them in using the mean, the median, or any other value.
 - drop the feature if most of the instances have a missing value.
 - drop the instance if you have several instances with missing values.
 - you can add an additional feature indicating whether the instance has a missing feature or not, and then use a value of 0 for the missing feature.

Data cleaning (II)

- Most ML methods require features that are numbers rather than categories (usually appearing as text).
- Also, if the feature is categorical, it is useful to use a different representation.
- In the previous bike rentals example, there were three categorical features
 - season that can take four categories.
 - holiday and functioning day, each taking two categories.

Data cleaning (III)

- For example, the feature season takes values autumn, winter, spring and summer.
- The way to handle this feature is to use a representation known as *one-hot encoding* to obtain a higher-dimensional binary representation for each value,

$$\text{autumn} = [1, 0, 0, 0]$$

$$\text{winter} = [0, 1, 0, 0]$$

$$\text{spring} = [0, 0, 1, 0]$$

$$\text{summer} = [0, 0, 0, 1]$$

- The values of the feature season do not have a natural order, therefore one should not map these values to numbers like 1 for autumn, 2 for winter, 3 for spring and 4 for summer.
- The ML method will try to find regularities within these ordered values even though they do not exist.

Feature selection and feature engineering

- You have the option to remove features that are uninformative.
- You have the option to discretise a continuous feature (e.g. binning).
- You can also create new features from the ones you have available.
- For example, instead of using feature x , you can use $\log(x)$, \sqrt{x} , x^2 , etc.

Feature scaling (I)

- Several ML methods do not perform well when the input features have very different scales.
- In the rental bikes example, the variable humidity is in the range 0 to 100, whereas wind speed is in the range 0 to 8.
- Two ways to get all features to have the same scale are *normalisation* (or min-max scaling) and *standardisation* (or z-score normalisation).

Feature scaling (II)

- In normalisation, we map the range of values that a feature takes to the range $[-1, 1]$ or $[0, 1]$.
- The normalisation formula is given as

$$\bar{x}_j = \frac{x_j - \min x_j}{\max x_j - \min x_j},$$

where $\min x_j$ and $\max x_j$ are the minimum and maximum values for the feature in the training set.

- In standardisation, the features are scaled so that they have mean zero and standard deviation equal to one,

$$\hat{x}_j = \frac{x_j - \mu_j}{\sigma_j},$$

where μ_j and σ_j are the mean and standard deviation of the feature x_j .

Feature scaling (III)

- Which scaling to use?
- If the dataset is not big and there is time, one can try both and see which one performs better in the validation set.
- If there is no time, as a rule of thumb:
 - unsupervised learning algorithms usually benefit more from standardisation than normalisation.
 - standardisation is preferred if the feature has already a distribution close to a Gaussian.
 - if the feature has outliers standardisation is preferred since normalisation will squeeze all the other values into a small range.
 - normalisation is preferred in all the other cases.

Contents

Frame the problem and look at the big picture

Get the data

Explore the data

Prepare the data

Shortlist models and then fine-tune them

Present your solution

Shortlist promising models

- ❑ If the data is big, use a smaller subset of the data to try different models in a reasonable time.
- ❑ Try different models from several categories (e.g. linear, non-linear, probabilistic, non-probabilistic).
- ❑ This can be easily done with scikit-learn.
- ❑ Measure and compare the performance of the ML models you used.
 - make sure you are using the exact same data in the training and the validation sets.
 - when using cross-validation, compute the mean and standard deviation of the performance measure over the k -folds for each model.
- ❑ shortlist two to three models that look promising.

Fine-tune the system (I)

- ❑ It is a good idea to use as much data as you can at this stage.
- ❑ Fine-tune the hyperparameters of your model using cross-validation.
 - data transformations are hyperparameters, e.g. should I input missing data with zero or with the mean?
 - use random search or grid search to explore hyperparameters
 - if the training takes long, use Auto-ML to fine-tune the hyperparameters.

Fine-tune the system (II)

- ❑ Once you are confident about your final model, usually the one that performs better on average on the validation data:
 - merge the training and the validation datasets
 - fit the model again using the set of hyperparameters you found before.
- ❑ Finally, use the test data to assess the generalisation error of your ML system.
 - Do not try to change your model based on the performance on the test data, you might start overfitting your model to the data.

Contents

Frame the problem and look at the big picture

Get the data

Explore the data

Prepare the data

Shortlist models and then fine-tune them

Present your solution

Present your solution

- Document what you have done.
- Create a nice presentation.
- Explain why your solution achieves the business objective.
- Bring up interesting points you noticed in the process.